

```
= What is RAG?
```

```
:order: 1
```

```
:type: lesson
```

```
[.slide]
```

```
-- Retrieval-Augmented Generation
```

Retrieval-Augmented Generation (RAG) is an approach that enhances the responses of LLMs by providing them with relevant, up-to-date information retrieved from external sources.

RAG helps generate more accurate and tailored answers, especially when the required information is not present in the model's training data.

```
[.slide.col-2]
```

```
-- The Retrieval-Augmented Generation (RAG) Process
```

```
[.col]
```

```
====
```

The RAG process typically involves three main steps:

- . **Understanding the User Query**

```
+
```

The system first interprets the user's input or question to determine what information is needed.

- . **Information Retrieval**

```
+
```

A `_retriever_` searches external data sources (such as documents, databases, or knowledge graphs) to find relevant information based on the user's query.

- . **Response Generation**

```
+
```

The retrieved information is inserted into the prompt, and the language model uses this context to generate a more accurate and relevant response.

```
====
```

```
[.col]
```

image::images/llm-rag-process.svg["A diagram showing the RAG process. A question from a user is sent to a retriever, which searches for relevant information. The retrieved information is then combined with the original question and sent to a language model, which generates a response."]

```
[.slide.col-2.discrete]
```

```
-- The RAG Process
```

```
[.col]
```

```
====
```

RAG systems can provide responses that are both contextually aware and grounded in real, up-to-date information.

If building a chatbot for a news agency, you could use RAG to pull real-time headlines or results from a news API.

When a user asks, "What's the latest news on the Olympics?", the chatbot, can provide a current headline or summary from the most recent articles, ensuring the response is timely and accurate.

[NOTE]

.Grounding

=====

The process of providing context to an LLM to improve the accuracy of its responses and reduce the likelihood of hallucinations is known as Grounding.

=====

=====

[.col]

image::images/llm-news-agency.svg["A news agency chatbot, showing the user asking a question, the chatbot grounding the question with a news API, and the chatbot responding with the latest news."]

[.slide]

== Retrievers

The retriever is a key component of the RAG process. A retriever is responsible for searching and retrieving relevant information from external data sources based on the user's query.

A retriever typically takes an *unstructured input* (like a question or prompt) and searches for structured data that can provide context or answers.

Neo4j support various methods for building retrievers, including:

- * Full-text search
- * Vector search
- * Text to Cypher

You will explore these methods in the rest of the course.

[.slide]

== Data sources

The data sources used in the RAG process can vary widely, depending on the application and the type of information needed. Common data sources include:

* **Documents**

+

Textual data sources, such as articles, reports, or manuals, that can be searched for relevant information.

* **APIs**

+

External services that can provide real-time data or specific information based on user queries.

```
* **Knowledge Graphs**  
+  
Graph-based representations of information that can provide context and  
relationships between entities.  
  
[ .slide.col-2.discrete]  
== Data sources  
  
[ .col]  
=====  
The news agency chatbot could use the following data sources:  
  
* A news API to retrieve the latest articles or headlines.  
* A knowledge graph to understand the relationships between different news  
topics, such as how they relate to each other or their historical context.  
This would help the chatbot provide more in-depth and contextual  
responses.  
* A document database to store and retrieve articles, reports, or other  
textual data that can be used to answer user queries.  
=====  
  
[ .col]  
image::images/llm-news-agency-knowledge-graph.svg["A news agency chatbot,  
showing the user asking a question, the chatbot grounding with the  
addition of data from a knowledge graph"]  
  
[ .transcript-only]  
=====  
[TIP]  
=====  
You will learn more about knowledge graphs and their construction in the  
next module.  
=====  
=====  
  
[ .quiz]  
== Check Your Understanding  
  
include::questions/1-rag.adoc[leveloffset=+1]  
  
[ .summary]  
== Lesson Summary  
  
In this lesson, you learned about RAG, which combines understanding user  
queries, retrieving relevant information, and generating a response using  
the retrieved information.  
  
In the next lesson, you will explore how embeddings and vector search can  
help retrieve relevant information.
```